

# Programming LSQUIC

Dmitri Tikhonov / LiteSpeed Technologies

Netdev 0x14

# Presentation outline

- Introduction
- LSQUIC history, features, and architecture
- Using API
  - Objects: engines, connections, streams
  - Sending and receiving packets
  - Instantiation and callbacks
- Bonus section, time permitting

## Example Program: tut.c

- You can follow along
- Snippets of `tut.c` are used as examples
- Compile the code while I talk about architecture

```
git clone https://github.com/litespeedtech/lsquic-tutorial
# This pulls in LSQUIC and BoringSSL
git submodule update --init --recursive
# Tested on Ubuntu
cmake .
make
```

# What is QUIC (keep it short)

- Why: Ossification
- Began as HTTP/2 over UDP
- Who: First Google, now everyone:
  - LiteSpeed Technologies, Microsoft, Apple, Facebook, Akamai, and others
- General-purpose transport protocol
- Killer feature: HTTP/3 (end of 2020)
- Experimental: datagrams, multipath, DNS-over-QUIC, NETCONF-over-QUIC, and so on
- Future: bright

# Which QUIC do you mean?

- Google QUIC vs IETF QUIC
- gQUIC? iQUIC? We all QUIC for yogurt?
- Google QUIC is on the way out
- In the slides that follow, “QUIC” means “IETF QUIC”

# Introducing LSQUIC

- Began as proprietary, then open-sourced
- Vanilla C
- Minimal dependencies
- Performance
- Robustness
- Flexibility

- *2016*. Goal: add Google QUIC support to LiteSpeed Web Server.
- *2017*. gQUIC support is shipped (Q035); LSQUIC is released on GitHub (client only).
- *2018*. IETF QUIC work begins. LSQUIC is the first functional HTTP/3 server.
- *2019*. LSQUIC 2.0 is released on GitHub (including the server bits). HTTP/3 support is shipped.
- *2020*. HTTP/3 is an RFC. (One can hope.)

- Latest (Draft-27, Draft-28, and Draft-29) IETF QUIC and HTTP/3 support, including
  - ECN, spin bits, path migration, NAT rebinding
  - Push promises, key updates
  - Several experimental extensions:
    - loss bits
    - timestamps
    - delayed ACKs
    - QUIC grease bit
- Google QUIC versions Q043, Q046, and Q050 (what Chrome currently uses)
- Many, many knobs to play with



- Bring your own event loop
- Bring your own networking
- Bring your own TLS context
- Scalable connection management

- Engine
- Connection
- Stream

- Client mode *or* server mode
- If need both, instantiate two
- HTTP mode

- Single library for both QUIC and HTTP/3
- Hide HTTP logic: control streams, header compression, framing
- Identical interface for gQUIC and HTTP/3
- Historical or strategic?
- Optimization: write-through

- Created, managed, and destroyed by engine
- Client initiates connection; object created before handshake is successful.
- Server: by the time user code gets the connection object, the handshake has already been completed
- Can have many streams during lifetime

- Belongs to a connection
- Bidirectional
- Usually corresponds to request/response exchange – depending on the application protocol
- API tries to mimic socket
  - But one can only take it so far

A single include file, contains all the necessary LSQUIC declarations:

```
#include "lsquic.h"
```

It pulls in auxiliary `lsquic_types.h`

# Library initialization

Before the first engine object is instantiated, the library must be initialized.

```
/* Example from tut.c */
if (0 != lsquic_global_init(LSQUIC_GLOBAL_SERVER
                           | LSQUIC_GLOBAL_CLIENT))
{
    fprintf(stderr, "global initialization failed\n");
    exit(EXIT_FAILURE);
}
```

This will initialize the crypto library, gQUIC server certificate cache, and, depending on the platform, monotonic timers.

If you plan to instantiate engines only in a single mode, client or server, you can omit the appropriate flag.



# Introducing `tut.c`

- Program to illustrate LSQUIC API use
- Contains both client and server code
- Echo service: client sends a line of text to server, the server returns the line, reversed
- Several examples that follow are excerpts from `tut.c`

# Running tutorial program

- Peruse online help: use the `-h` flag
- Running client or server: the server takes `-c` and `-k` arguments

# Server:

```
sh$ ./tut -c mycert-cert.pem -k mycert-key.pem ::0 12345
```

# Client:

```
sh$ ./tut ::1 12345 -L debug -f client.log
```

Hello!

!olleH

^D

# Engine constructor

- Server or client
- HTTP mode

```
/* Return a new engine instance.  
* `flags' is bitmask of LSENG_SERVER and LSENG_HTTP.  
* `api' is required.  
*/  
lsquic_engine_t *  
lsquic_engine_new (unsigned flags,  
                  const struct lsquic_engine_api *api);
```

# Specifying engine callbacks

- Pass pointer to struct `lsquic_engine_api`

```
/* Minimal configuration */  
struct lsquic_engine_api engine_api = {  
    .ea_packets_out      = send_packets_out,  
    .ea_packets_out_ctx  = sender_ctx,  
    .ea_stream_if        = &stream_callbacks,  
    .ea_stream_if_ctx    = &some_context,  
    .ea_get_ssl_ctx      = get_ssl_ctx, /* Server only */  
};
```

## Excerpt from tut.c

```
/* Initialize callbacks */
memset(&eapi, 0, sizeof(eapi));
eapi.ea_packets_out = tut_packets_out;
eapi.ea_packets_out_ctx = &tut;
eapi.ea_stream_if    = tut.tut_flags & TUT_SERVER
                      ? &tut_server_callbacks : &tut_client_callbacks;
eapi.ea_stream_if_ctx = &tut;
eapi.ea_get_ssl_ctx   = tut_get_ssl_ctx;
```

# Packets in

- Single function to feed packets to engine instance
- Specify: datagram, peer and local addresses, ECN value
- `peer_ctx` is associated with peer address: it is passed to `send_packets` function.

```
/* 0: processed by real connection  
 * 1: handled  
 *-1: error: invalid arguments, malloc failure  
 */  
int  
lsquic_engine_packet_in (lsquic_engine_t *,  
    const unsigned char *udp_payload, size_t sz,  
    const struct sockaddr *sa_local,  
    const struct sockaddr *sa_peer,  
    void *peer_ctx, int ecn);
```

# Why specify local address

- Becomes source address on outgoing packets
  - Important in multihomed configuration
- Path change detection
  - QUIC sends special frames to validate path

# Packets out

- Callback
- Called during connection processing (explicit call by user)

```
/* Returns number of packets successfully sent out or -1 on error.
 *
 * If not all packets could be sent, call
 * lsquic_engine_send_unsent_packets() when can send again.
 */
typedef int (*lsquic_packets_out_f)(
    void                *packets_out_ctx,
    const struct lsquic_out_spec *out_spec,
    unsigned            n_packets_out
);
```



## When an error occurs

- `errno` is examined
- `EAGAIN` (or `EWOULDBLOCK`) means retry later; engine enters the “can’t send packets” mode
- Call `lsquic_engine_send_unsent_packets()` when sending is possible again
- `EMSGSIZE` means the packet is too large. This happens when MTU probes are sent. The engine retries sending without the offending packet.
- Other `errno` values cause immediate termination of corresponding connection.

# Outgoing packet specification

- Why `iovec`?
  - UDP datagram can contain more than one QUIC packet
  - Packet coalescing

```
struct lsquic_out_spec
{
    struct iovec          *iov;
    size_t                iovlen;
    const struct sockaddr *local_sa;
    const struct sockaddr *dest_sa;
    void                  *peer_ctx;
    int                   ecn; /* 0 - 3; see RFC 3168 */
};
```

## Packets out example

```
static int
my_packets_out (void *ctx, const struct lsquic_out_spec *specs,
                unsigned n_specs) {
    struct msghdr msg;  memset(&msg, 0, sizeof(msg));
    unsigned n;
    for (n = 0; n < n_specs; ++n) {
        msg.msg_name      = (void *) specs[n].dest_sa;
        msg.msg_namelen   = sizeof(struct sockaddr_in);
        msg.msg_iov       = specs[n].iov;
        msg.msg_iovlen    = specs[n].iovlen;
        if (sendmsg((int) specs[n].peer_ctx, &msg, 0) < 0)
            break;
    }
    return (int) n;
}
```

# When to process connections

- Connections are either tickable immediately or at some future point
- Future point may be a retransmission, idle, or some other alarm
- Only exception when idle timeout is disabled (on by default)
- Engine knows when to process connections next

```
/* Returns true if there are connections to be processed, in  
 * which case `diff' is set to microseconds from current time.  
 */  
int  
lsquic_engine_earliest_adv_tick (lsquic_engine_t *, int *diff);
```

## Example with event loop

```
/* Abbreviated, see full version in tut.c */
void tut_process_conns (struct tut *tut) {
    ev_tstamp timeout;
    int diff;
    ev_timer_stop();
    lsquic_engine_process_conns(engine);
    if (lsquic_engine_earliest_adv_tick(engine, &diff) {
        if (diff > 0)
            timeout = (ev_tstamp) diff / 1000000;    /* To seconds */
        else
            timeout = 0.;
        ev_timer_init(timeout)
        ev_timer_start();
    }
}
```

# Tickable connection

- There are incoming packets
- A stream is both readable by the user code and the user code wants to read from it
- A stream is both writeable by the user code and the user code wants to write to it
- User has written to stream outside of `on_write()` callbacks (that is allowed) and now there are packets ready to be sent
- A control frame needs to be sent out
- A stream needs to be serviced or created

# Required engine callbacks

- Callback to send packets
- Connection and stream callbacks
  - `on_new_conn()`, `on_read()`, and so on
- Callback to get default TLS context (server only)

# Optional callbacks

- Certificate lookup by SNI (server only)
- Outgoing packet memory allocation
- Connection ID lifecycle: new, live, and old CIDs
- Shared memory hash
- HTTP header set processing



# Stream and connection callbacks

- Specified in `struct lsquic_stream_if`
- Mandatory callbacks:
  - `on_new_conn()` - new connection is created
  - `on_conn_closed()`
  - `on_new_stream()`
  - `on_read()`
  - `on_write()`
  - `on_close()`
- Optional callbacks:
  - `on_goaway_received()`
  - `on_new_token()` (client only)
  - `on_hsk_done()` (client only)
  - `on_sess_resume_info()` (client only)

## On new connection

- Server: handshake successful; client: object created
- Chance to create custom per-connection context

```
/* Return pointer to per-connection context. OK to return NULL. */
static lsquic_conn_ctx_t *
my_on_new_conn (void *ea_stream_if_ctx, lsquic_conn_t *conn)
{
    struct some_context *ctx = ea_stream_if_ctx;
    struct my_conn_ctx *my_ctx = my_ctx_new(ctx);
    if (ctx->is_client)
        /* Need a stream to send request */
        lsquic_conn_make_stream(conn);
    return (void *) my_ctx;
}
```

## On new stream

- Depending on situation, register interest in reading or writing
  - Or just read or write
- Chance to create per-stream context

```
/* Return pointer to per-connection context. OK to return NULL. */
static lsquic_stream_ctx_t *
my_on_new_stream (void *ea_stream_if_ctx, lsquic_stream_t *stream) {
    struct some_context *ctx = ea_stream_if_ctx;
    /* Associate some data with this stream: */
    struct my_stream_ctx *stream_ctx
        = my_stream_ctx_new(ea_stream_if_ctx);
    stream_ctx->stream = stream;
    if (ctx->is_client)
        lsquic_stream_wantwrite(stream, 1);
    return (void *) stream_ctx;
}
```

# On read

- Read data – or collect error

```
static void
my_on_read (lsquic_stream_t *stream, lsquic_stream_ctx_t *h) {
    struct my_stream_ctx *my_stream_ctx = (void *) h;
    unsigned char buf[BUFSZ];

    ssize_t nr = lsquic_stream_read(stream, buf, sizeof(buf));
    /* Do something with the data.... */
    if (nr == 0) /* EOF */ {
        lsquic_stream_shutdown(stream, 0);
        lsquic_stream_wantwrite(stream, 1); /* Want to reply */
    }
}
```

- If called, you should be able to write *some* bytes

```
static void
my_on_write (lsquic_stream_t *stream, lsquic_stream_ctx_t *h) {
    struct my_stream_ctx *my_stream_ctx = (void *) h;
    ssize_t nw = lsquic_stream_write(stream,
        my_stream_ctx->resp, my_stream_ctx->resp_sz);
    if (nw == my_stream_ctx->resp_sz)
        lsquic_stream_close(stream);
}
```

# On stream close

- After this, stream will be destroyed, drop all pointers to it

```
/* Made-up example */
static void
my_on_close (lsquic_stream_t *stream, lsquic_stream_ctx_t *h) {
    lsquic_conn_t *conn = lsquic_stream_conn(stream);
    struct my_conn_ctx *my_ctx = lsquic_conn_get_ctx(conn);
    if (!has_more_reqs_to_send(my_ctx)) /* For example */
        lsquic_conn_close(conn);
    free(h);
}
```

## On stream close in tut.c

```
static void
tut_server_on_close (lsquic_stream_t *stream,
                    lsquic_stream_ctx_t *h)
{
    struct tut_server_stream_ctx *const tssc = (void *) h;
    free(tssc);
    LOG("stream closed");
}
```

## On connection close

- After this, connection will be destroyed, drop all pointers to it

```
static void
my_on_conn_closed (lsquic_conn_t *conn) {
    struct my_conn_ctx *my_ctx = lsquic_conn_get_ctx(conn);
    struct some_context *ctx = my_ctx->some_context;

    --ctx->n_conns;
    if (0 == ctx->n_conn && (ctx->flags & CLOSING))
        exit_event_loop(ctx);

    free(my_ctx);
}
```



## On connection close in tut.c

- This deletes last event from the event loop – the event loop exits and the program terminates
- Note that the same thing happens if handshake fails

```
static void
tut_client_on_conn_closed (struct lsquic_conn *conn)
{
    struct tut *const tut = (void *) lsquic_conn_get_ctx(conn);

    LOG("client connection closed -- stop reading from socket");
    ev_io_stop(tut->tut_loop, &tut->tut_sock_w);
}
```

# More about streams

- When writing to a stream, data is placed directly into packets
  - Except when it isn't
- Writes smaller than packet size are buffered
- Inside `on_read` and `on_write` callbacks, reading and writing succeeds
- Unless stream is reset, in which case reading from stream returns -1
  - This is done so that user can collect error

# More stream functions

```
/* Flush any buffered data. This triggers packetizing even a single  
 * byte into a separate frame.  
 */
```

```
int  
lsquic_stream_flush (lsquic_stream_t *);
```

```
/* Possible values for how are 0, 1, and 2. See shutdown(2). */
```

```
int  
lsquic_stream_shutdown (lsquic_stream_t *, int how);
```

```
int  
lsquic_stream_close (lsquic_stream_t *);
```

# Stream return values and error codes

- Reading and writing interface modeled on `read(2)` and `write(2)`
- Including the use of `errno`
  - If no data to read, error is `EWOULDBLOCK`
  - If stream is closed, error is `EBADF`
  - If stream is reset, error is `ECONNRESET`
  - After this, error codes fit only if you squint very hard
    - `EINVAL` - argument to `shutdown` is not 0, 1, or 2
    - `EILSEQ` - cannot send HTTP payload before headers
    - `EBADMSG` - sending HTTP headers is not allowed (several reasons)
- When `lsquic_stream_read()` returns 0, it means EOF
- `lsquic_stream_write()` returns 0 when flow control or congestion control limit is reached.

# More ways to read and write

- Scatter/gather
- Similar to `readv(2)` and `writew(2)`

`ssize_t`

```
lsquic_stream_readv (lsquic_stream_t *, const struct iovec *,  
                                int iovcnt);
```

`ssize_t`

```
lsquic_stream_writew (lsquic_stream_t *, const struct iovec *,  
                                int count);
```

# Read using a callback

- Use for zero-copy stream processing
- `lsquic_stream_read()` and `lsquic_stream_readv()` are just wrappers
- Callback returns number of bytes processed
  - Pointer to user-supplied context;
  - Pointer to the data;
  - Data size (can be zero); and
  - Indicator whether the FIN follows the data.
- If callback returns 0 or value smaller than `len`, reading stops

`ssize_t`

```
lsquic_stream_readf (lsquic_stream_t *,  
    size_t (*readf)(void *ctx, const unsigned char *, size_t len, int fin),  
    void *ctx);
```

## Stream read: copy data

- The classic way to read
- Easy, but it copies data

```
static void    /* v0 is the default.  On command line: -b 0 */
tut_client_on_read_v0 (lsquic_stream_t *stream, lsquic_stream_ctx_t *h)
{
    struct tut *tut = (struct tut *) h;
    unsigned char buf[3];
    ssize_t nread = lsquic_stream_read(stream, buf, sizeof(buf));
    if (nread > 0)
    {
        fwrite(buf, 1, nread, stdout);
        fflush(stdout);
    }
    /* --- 8< --- snip --- 8< --- */
}
```

## Stream read take 2: use callback

```
static void    /* On command line: -b 1 */
tut_client_on_read_v1 (lsquic_stream_t *stream, lsquic_stream_ctx_t *h)
{
    struct tut *tut = (struct tut *) h;
    size_t nread = lsquic_stream_readf(stream, tut_client_readf_v1, NULL);
    if (nread == 0)
    {
        LOG("read to end-of-stream: close and read from stdin again");
        lsquic_stream_shutdown(stream, 0);
        ev_io_start(tut->tut_loop, &tut->tut_u.c.stdin_w);
    }
    /* --- 8< --- snip --- 8< --- */
}
```



## Stream read take 2: the callback itself

- `len` may be arbitrary (and larger than packet size)
- Return value smaller than `len` to stop callback

```
static size_t
tut_client_readf_v1 (void *ctx, const unsigned char *data,
                    size_t len, int fin)
{
    if (len)
    {
        fwrite(data, 1, len, stdout);
        fflush(stdout);
    }
    return len;
}
```

## Stream read take 3: use FIN

- Save one `on_read()` call

```
struct client_read_v2_ctx { struct tut *tut; lsquic_stream_t *stream; };

static void /* On command line: -b 2 */
tut_client_on_read_v2 (lsquic_stream_t *stream,
                      lsquic_stream_ctx_t *h) {
    struct tut *tut = (struct tut *) h;
    struct client_read_v2_ctx v2ctx = { tut, stream, };
    ssize_t nread = lsquic_stream_readf(stream, tut_client_readf_v2,
                                         &v2ctx);

    if (nread < 0)
        /* ERROR */
}
```

## Stream read take 3: the callback itself

```
static size_t
tut_client_readf_v2 (void *ctx, const unsigned char *data,
                    size_t len, int fin) {
    struct client_read_v2_ctx *v2ctx = ctx;
    if (len)
        fwrite(data, 1, len, stdout);
    if (fin)
    {
        fflush(stdout);
        LOG("read to end-of-stream: close and read from stdin again");
        lsquic_stream_shutdown(v2ctx->stream, 0);
        ev_io_start(v2ctx->tut->tut_loop, &v2ctx->tut->tut_u.c.stdin_w);
    }
    return len;
}
```

## Stream write take 1

```
static void    /* v0 is the default. On command line: -w 0 */
tut_server_on_write_v0 (lsquic_stream_t *stream, lsquic_stream_ctx_t *h)
{
    struct tut_server_stream_ctx *const tssc = (void *) h;
    ssize_t nw = lsquic_stream_write(stream,
        tssc->tssc_buf + tssc->tssc_off, tssc->tssc_sz - tssc->tssc_off);
    if (nw > 0)
    {
        tssc->tssc_off += nw;
        if (tssc->tssc_off == tssc->tssc_sz)
            lsquic_stream_close(stream);
    }
    /* --- 8< --- snip --- 8< --- */
}
```

# Write using callbacks

```
struct lsquic_reader {  
    /* Return number of bytes written to buf */  
    size_t (*lsqr_read) (void *lsqr_ctx, void *buf, size_t count);  
    /* Return number of bytes remaining in the reader. */  
    size_t (*lsqr_size) (void *lsqr_ctx);  
    void *lsqr_ctx;  
};  
  
/* Return number of bytes written or -1 on error. */  
ssize_t  
lsquic_stream_writef (lsquic_stream_t *, struct lsquic_reader *);
```

## Stream write take 2

- Useful when reading from external data source, such as file descriptor
- Write data directly into stream frame

```
static void    /* On command line: -w 1 */
tut_server_on_write_v1 (lsquic_stream_t *stream, lsquic_stream_ctx_t *h)
{
    struct tut_server_stream_ctx *const tssc = (void *) h;
    struct lsquic_reader reader = { tssc_read, tssc_size, tssc, };
    ssize_t nw = lsquic_stream_writef(stream, &reader);
    if (nw > 0 && tssc->tssc_off == tssc->tssc_sz)
        lsquic_stream_close(stream);
    /* --- 8< --- snip --- 8< --- */
}
```

# Reader size callback

- Return number of bytes remaining

```
static size_t  
tssc_size (void *ctx)  
{  
    struct tut_server_stream_ctx *tssc = ctx;  
    return tssc->tssc_sz - tssc->tssc_off;  
}
```

# Reader read callback

- `count` is calculated using `tssc_size()`
- If larger than amount of remaining data, may indicate truncation

```
static size_t
tssc_read (void *ctx, void *buf, size_t count)
{
    struct tut_server_stream_ctx *tssc = ctx;

    if (count > tssc->tssc_sz - tssc->tssc_off)
        count = tssc->tssc_sz - tssc->tssc_off;
    memcpy(buf, tssc->tssc_buf + tssc->tssc_off, count);
    tssc->tssc_off += count;
    return count;
}
```



## Client: making connection

```
lsquic_conn_t *  
lsquic_engine_connect (lsquic_engine_t *,  
    enum lsquic_version, /* Set to N_LSQVER for default */  
    const struct sockaddr *local_sa,  
    const struct sockaddr *peer_sa,  
    void *peer_ctx,  
    lsquic_conn_ctx_t *conn_ctx,  
    const char *hostname, /* Used for SNI */  
    unsigned short base_plpmtu, /* 0 means default */  
    const unsigned char *sess_resume, size_t sess_resume_len,  
    const unsigned char *token, size_t token_sz);
```

## Excerpt from tut.c

```
tut.tut_u.c.conn = lsquic_engine_connect(
    tut.tut_engine, N_LSQVER,
    (struct sockaddr *) &tut.tut_local_sas, &addr.sa,
    (void *) (uintptr_t) tut.tut_sock_fd, /* Peer ctx */
    NULL, NULL, 0, NULL, 0, NULL, 0);
if (!tut.tut_u.c.conn)
{
    LOG("cannot create connection");
    exit(EXIT_FAILURE);
}
tut_process_conns(&tut);
```

# Specifying QUIC version

- There is no version negotiation version in QUIC... yet
- When passed to `lsquic_engine_connect`, `N_LSQVER` means “let the engine pick the version”
  - The engine picks the highest it supports, so that’s a good way to go

```
enum lsquic_version {  
    LSQVER_043, LSQVER_046, LSQVER_050,      /* Google QUIC */  
    LSQVER_ID27, LSQVER_ID28, LSQVER_ID29,    /* IETF QUIC */  
    /* ...some special entries skipped */      N_LSQVER  
};
```

```
/* This allows list of versions to be specified as bitmask: */  
es_versions = (1 << LSQVER_ID28) | (1 << LSQVER_ID29);
```

## Server: additional callbacks

- SSL context and certificate callbacks

```
typedef struct ssl_ctx_st * (*lsquic_lookup_cert_f)(  
    void *lsquic_cert_lookup_ctx, const struct sockaddr *local,  
    const char *sni);  
  
struct lsquic_engine_api {  
    lsquic_lookup_cert_f    ea_lookup_cert;  
    void                    *ea_cert_lu_ctx;  
    struct ssl_ctx_st *    (*ea_get_ssl_ctx)(void *peer_ctx);  
    /* (Other members of the struct are not shown) */  
};
```

- Engine API's `ea_settings` may be pointed to settings.
- `struct lsquic_engine_settings`.
- There are *many* settings (over 50).
- Do use `lsquic_engine_init_settings()`.

# Settings helper functions

```
/* Initialize 'settings' to default values */
void
lsquic_engine_init_settings (struct lsquic_engine_settings *,
    /* Bitmask of LSENG_SERVER and LSENG_HTTP */
    unsigned lsquic_engine_flags);

/* Check settings for errors, return 0 on success, -1 on failure. */
int
lsquic_engine_check_settings (const struct lsquic_engine_settings *,
    unsigned lsquic_engine_flags,
    /* Optional, can be NULL: */
    char *err_buf, size_t err_buf_sz);
```

## Settings example in tut.c 1/2

```
case 'o':    /* For example: -o version=h3-27 -o cc_algo=2 */
    if (!settings_initialized) {
        lsquic_engine_init_settings(&settings,
                                     cert_file || key_file ? LSENG_SERVER : 0);
        settings_initialized = 1;
    }
    /* ... */

else if (0 == strncmp(optarg, "cc_algo=", val - optarg))
    settings.es_cc_algo = atoi(val);
```

## Settings example in tut.c 2/2

```
/* Check settings */
if (0 != lsquic_engine_check_settings(&settings,
                                     tut.tut_flags & TUT_SERVER ? LSENG_SERVER : 0,
                                     errbuf, sizeof(errbuf)))
{
    LOG("invalid settings: %s", errbuf);
    exit(EXIT_FAILURE);
}

/* ... */
eapi.ea_settings = &settings;
```



# Logging mechanism

- By default, log messages are thrown away

```
struct lsquic_logger_if {  
    int (*log_buf)(void *logger_ctx, const char *buf, size_t len);  
};
```

```
enum lsquic_logger_timestamp_style { LLTS_NONE, LLTS_HHMMSSMS,  
    LLTS_YYYYMMDD_HHMMSSMS, LLTS_CHROMELIKE, LLTS_HHMMSSUS,  
    LLTS_YYYYMMDD_HHMMSSUS, N_LLTS };
```

```
void lsquic_logger_init(const struct lsquic_logger_if *,  
    void *logger_ctx, enum lsquic_logger_timestamp_style);
```

# Logging levels and modules

- Eight log levels
  - Debug, info, notice, warning, error, alert, emerg, crit.
  - Only first five are used; warning is the default
- Many modules
  - Event, engine, stream, connection, bbr, and many more
  - Refer to documentation

```
/* Set log level for all modules */  
int lsquic_set_log_level (const char *log_level);
```

```
/* Set log level per module "event=debug" */  
int lsquic_logger_lopt (const char *optarg);
```

## Logging in tut.c 1/2

- Check out -f, -l, and -L flags
- E.g. -f log.file (goes to *stderr* by default)
- E.g. -l event=debug,stream=info or -L debug

```
static int
tut_log_buf (void *ctx, const char *buf, size_t len) {
    FILE *out = ctx;
    fwrite(buf, 1, len, out);
    fflush(out);
    return 0;
}
static const struct lsquic_logger_if logger_if = { tut_log_buf, };

lsquic_logger_init(&logger_if, s_log_fh, LLTS_HHMMSSUS);
```

## Logging in tut.c 2/2

```
case 'l':    /* e.g. -l event=debug,cubic=info */
    if (0 != lsquic_logger_lopt(optarg)) {
        fprintf(stderr, "error processing -l option\n");
        exit(EXIT_FAILURE);
    }
    break;
case 'L':    /* e.g. -L debug */
    if (0 != lsquic_set_log_level(optarg)) {
        fprintf(stderr, "error processing -L option\n");
        exit(EXIT_FAILURE);
    }
    break;
```

- Wireshark supports IETF QUIC
- Need version 3.3 for Internet-Draft 29 support
- Export TLS secrets using `ea_keylog_if`
- Example: `-G` option in `tut`:

```
if (key_log_dir) {  
    eapi.ea_keylog_if = &keylog_if;  
    eapi.ea_keylog_ctx = (void *) key_log_dir;  
}
```

# LSQUIC API for key logging

```
/* Secrets are logged per connection. Interface to open file (handle),  
 * log lines, and close file.  
 */
```

```
struct lsquic_keylog_if {  
    void * (*kli_open) (void *keylog_ctx, lsquic_conn_t *);  
    void (*kli_log_line) (void *handle, const char *line);  
    void (*kli_close) (void *handle);  
};
```

```
struct lsquic_engine_api {  
    /* --- 8< --- snip --- 8< --- */  
    const struct lsquic_keylog_if *ea_keylog_if;  
    void *ea_keylog_ctx;  
};
```

# Key logging in tut.c

```
static struct lsquic_keylog_if keylog_if = {
    .kli_open      = keylog_open,
    .kli_log_line   = keylog_log_line,
    .kli_close      = keylog_close,
};

static void
keylog_log_line (void *handle, const char *line)
{
    fputs(line, handle);
    fputs("\n", handle);
    fflush(handle);
}
```

# Wireshark screenshot

No.	Time	Source	Destination	Protocol	Length	Info
16	2.570135	127.0.0.1	127.0.0.1	QUIC	83	Protected Payload (KP0), DCID=e77ce754b2d32f6c, PKN: 7, ACK
17	2.924101	127.0.0.1	127.0.0.1	QUIC	78	Protected Payload (KP0), DCID=e77ce754b2d32f6c, PKN: 8, STREAM(0)
18	2.925751	127.0.0.1	127.0.0.1	QUIC	90	Protected Payload (KP0), DCID=be3ac0381d9049e8, PKN: 11, ACK, STREAM(...
19	2.927297	127.0.0.1	127.0.0.1	QUIC	81	Protected Payload (KP0), DCID=e77ce754b2d32f6c, PKN: 9, ACK
20	2.942133	127.0.0.1	127.0.0.1	QUIC	97	Protected Payload (KP0), DCID=be3ac0381d9049e8, PKN: 12, NCI
21	2.944590	127.0.0.1	127.0.0.1	QUIC	81	Protected Payload (KP0), DCID=e77ce754b2d32f6c, PKN: 10, ACK

## ACK

Frame Type: ACK (0x0000000000000002)

Largest Acknowledged: 8

ACK Delay: 46

ACK Range Count: 0

First ACK Range: 3

## TIME\_STAMP

Frame Type: TIME\_STAMP (0x00000000000002f5)

Time Stamp: 365556

## STREAM id=0 fin=1 off=0 len=7 uni=0

### Frame Type: STREAM (0x000000000000000b)

....1 = Fin: True

....1. = Len(gth): True

....0.. = Off(set): False

Stream ID: 0

Length: 7

Stream Data: 216f6c6c65480a

0000	02 08 2e 00 03 42 f5 80 05 93 f4 0b 00 07 21 6f	...B... ..!c
0010	6c 6c 65 48 0a	lleH.



# Connection ID

```
#define MAX_CID_LEN 20

typedef struct lsquic_cid
{
    uint_fast8_t    len;
    union {
        uint8_t     buf[MAX_CID_LEN];
        uint64_t     id;
    }               u_cid;
#define idbuf u_cid.buf
} lsquic_cid_t;

#define LSQUIC_CIDS_EQ(a, b) ((a)->len == 8 ? \
    (b)->len == 8 && (a)->u_cid.id == (b)->u_cid.id : \
    (a)->len == (b)->len && 0 == memcmp((a)->idbuf, (b)->idbuf, (a)->len))
```

# Get this-and-that API

```
const lsquic_cid_t *  
lsquic_conn_id (const lsquic_conn_t *);  
  
lsquic_conn_t *  
lsquic_stream_conn (const lsquic_stream_t *);  
  
lsquic_engine_t *  
lsquic_conn_get_engine (lsquic_conn_t *);  
  
int  
lsquic_conn_get_sockaddr (lsquic_conn_t *,  
                           const struct sockaddr **local, const struct sockaddr **peer);
```

# Stream priorities

- 1 through 256, where 1 is the highest priority
- Controls dispatch of `on_read()` and `on_write()` callbacks
- *Not* HTTP/3 priorities (those are coming later)

```
/* Set stream priority. Valid priority values are 1 through 256,  
 * inclusive. Lower value means higher priority.  
 */
```

```
int
```

```
lsquic_stream_set_priority (lsquic_stream_t *, unsigned priority);
```

```
/* Return current priority of the stream */
```

```
unsigned
```

```
lsquic_stream_priority (const lsquic_stream_t *);
```

## But wait, there is more!

- LSQUIC has more features that we did not cover
- Shared memory callbacks to support persistent crypto contexts and multi-process setups
- Callbacks to handle memory allocation for outgoing packets
- Callbacks to observe CID life cycle: new, live, and dead. Useful in multi-process setups
- Server certificate verification callbacks
- And more: auto-tuning; r/w callback no-progress detection; r/w callback dispatch: once vs loop; different congestion controls; optimistic ACK attack detection; stream fragmentation, reassembly, and commitment attack mitigation
- Refer to documentation and to more involved example programs in the LSQUIC distribution
- <https://lsquic.readthedocs.io/>

\_\_END\_\_

## Linux Wishlist

# GSO and ECN

- Rather, GSO *or* ECN
- Cannot have both

```
/* Have kernel split `packet_sz` bytes into packets... */
```

```
cmmsg->cmmsg_level = SOL_UDP;  
cmmsg->cmmsg_type = UDP_SEGMENT;  
cmmsg->cmmsg_len = CMSG_LEN(sizeof(uint16_t));  
*((uint16_t *) CMSG_DATA(cmmsg)) = packet_sz;
```

```
/* ... but how can one specify ECN for each? */
```

```
cmmsg->cmmsg_level = IPPROTO_IP;  
cmmsg->cmmsg_type = IP_TOS;  
cmmsg->cmmsg_len = CMSG_LEN(sizeof(tos));  
memcpy(CMSG_DATA(cmmsg), &tos, sizeof(tos));
```

# DPLPMTUD: Suppressing EMSGSIZE

ip(7) has this to say:

*It is possible to implement RFC 4821 MTU probing with SOCK\_DGRAM or SOCK\_RAW sockets by setting a value of IP\_PMTUDISC\_PROBE (available since Linux 2.6.22). This is also particularly useful for diagnostic tools such as tracepath(8) that wish to deliberately send probe packets larger than the observed Path MTU.*

```
/* Even with this setting, sendmsg(2) returns -1 with EMSGSIZE  
 * if datagram larger than local interface's MTU  
 */
```

```
on = IP_PMTUDISC_PROBE;
```

```
s = setsockopt(fd, IPPROTO_IP, IP_MTU_DISCOVER, &on, sizeof(on));
```

- Painful: have to handle EMSGSIZE specially and retry `sendmmsg(2)`

### HTTP/3



# HTTP/3 differences

- ALPN is required
  - This step is handled by the library
  - QUIC version I-D 29 corresponds to ALPN “h3-29” and so on
  - QUIC version 1 will correspond to “h3”
- SNI is required
  - Pass it to `lsquic_engine_connect()`
- Send headers before sending payload
  - Use `lsquic_stream_send_headers()`
- Optional: header callbacks via `ea_hsi_if` (HSI: header set interface)
  - If not specified, LSQUIC will pretend you are reading HTTP/1.x-like stream.
- Server must use `ea_lookup_cert` callback

## h3cli.c: a simple HTTP/3 client

- Specify hostname, port number (or service name), and path
- Default method is **GET**
- Websites to try:
  - [www.litespeedtech.com](http://www.litespeedtech.com)
  - [www.google.com](http://www.google.com)
  - [www.facebook.com](http://www.facebook.com)

```
./h3cli www.litespeedtech.com 443 / -M HEAD
```

## h3cli.c: use LSENG\_HTTP flag

```
/* The following three functions take the HTTP flag: */

lsquic_engine_init_settings(&settings, LSENG_HTTP);

if (0 != lsquic_engine_check_settings(&settings, LSENG_HTTP,
                                     errbuf, sizeof(errbuf)))
    /* error */ ;

h3cli.h3cli_engine = lsquic_engine_new(LSENG_HTTP, &eapi);
```

```
h3cli.h3cli_conn = lsquic_engine_connect(  
    h3cli.h3cli_engine, N_LSQVER,  
    (struct sockaddr *) &h3cli.h3cli_local_sas, &addr.sa,  
    (void *) (uintptr_t) h3cli.h3cli_sock_fd,  /* Peer ctx */  
    NULL,  
    h3cli.h3cli_hostname,  /* <== This becomes SNI in ClientHello */  
    0, NULL, 0, NULL, 0);
```

## h3cli.c: send requests

```
static void h3cli_client_on_write (struct lsquic_stream *stream,
                                   lsquic_stream_ctx_t *h) {

    struct header_buf hbuf;
    struct lsxpack_header harray[5];
    struct lsquic_http_headers headers = { 5, harray, };

    h3cli_set_header(&harray[0], &hbuf, V(":method"), V("GET"));
    h3cli_set_header(&harray[1], &hbuf, V(":scheme"), V("https"));
    /* --- 8< --- snip --- 8< --- */

    if (0 == lsquic_stream_send_headers(stream, &headers, 0))
    {
        lsquic_stream_shutdown(stream, 1);
        lsquic_stream_wantread(stream, 1);
    }
}
```

## HTTP/3: configuration options

- What's a few more tunable parameters between friends?
- Push promises
- Compression
  - QPACK choices: table size; allow blocked streams?
- Priorities (coming as an extension)
- LSQUIC aims for sane defaults
- Change to fit your needs and performance goals; YMMV